VIKAS YADAV

Senior SRE / DevOps Engineer · Platform & Infrastructure · AWS · Terraform · Observability · MLOps

vikaskashi896@gmail.com · +91 72485 56138 (WhatsApp) · Hyderabad, India
linkedin.com/in/vikas-yadav896 · github.com/iamyadavvikas · Open to Bangalore / Pune / Hyderabad / Noida / Remote-India · Immediate joiner

SUMMARY

Platform & Infrastructure SRE / DevOps engineer with **4 years** at Shutterfly Inc. (consulting via Genpact), working on the Photos & Media platform. I own the Terraform code for our AWS setup across multiple accounts and **four** environments (dev, beta, stage, prod), with **100+** pull requests merged after peer review.

Day to day I work with AWS (ECS, EC2/ASG, S3, IAM, ELB/ALB, VPC, SSM, Route53), Terraform, Packer for golden AMIs, and Jenkins Job DSL. I built and run our Splunk logging platform (license master plus heavy forwarders) from scratch. I also handle ECS capacity, IAM least-privilege cleanups, network security hardening, infrastructure decommissioning, and a **24×7** on-call rotation.

I run the infrastructure that sits under our media and AI services — the image pipeline, the recommendation engine, and AI config — rather than training the models myself. MLOps and LLM-serving infrastructure (vLLM, KServe, MLflow) are an active learning track for me right now.

SKILLS

Infrastructure as Code: Terraform (advanced — multi-account, multi-environment, modules, plan/apply via PR), Packer (golden AMIs / immutable infrastructure), CloudFormation

Cloud & Linux: AWS (ECS, EC2 / Auto Scaling Groups & Launch Templates, S3, Lambda, SQS, IAM, ELB / ALB, VPC / Security Groups, SSM, Route53), Linux (RHEL, Ubuntu), Shell scripting

CI/CD & Automation: Jenkins (Job DSL, Pipeline DSL, Groovy), Git / GitHub Enterprise, infrastructure & jobs as code, deployment automation, environment provisioning, decommissioning automation

Containers & Orchestration: Docker, Amazon ECS (capacity & service scaling)

Observability & SRE: Splunk (built & operate the logging platform), SignalFX, VividCortex, CloudWatch, RUM / synthetics / APM consumption, SLOs / SLIs, incident response, RCAs, fast revert-based rollback

Security: IAM design & least-privilege cleanup, internal-only ELB lockdown, security groups / network segmentation, CVE / vulnerability exception management, bastion / jumphost access

Programming & Scripting: Python (advanced), Bash, Groovy, SQL, REST APIs

Familiar / actively learning: EKS, Helm, ArgoCD, Prometheus, Grafana, Ansible, GitHub Actions, OpenTelemetry, PagerDuty, vLLM, KServe, Triton, MLflow

EXPERIENCE

Site Reliability & DevOps Engineer (Platform & Infrastructure) | 08/2022 - Present
Shutterfly Inc. (via Genpact), Hyderabad - Photos & Media Platform

- Own the Terraform code for the Photos platform across multiple AWS accounts and **four** environments (dev, beta, stage, prod). **100+** merged pull requests, separate state per environment, shared modules, and a PR + `terraform plan` review before anything goes to production.
- Built our Splunk logging platform from scratch — license master and heavy forwarders — using Terraform Launch Templates, Auto Scaling Groups, and AMIs. This gives us a self-healing, code-managed logging layer the whole team uses for troubleshooting.
- Maintain a Packer golden-AMI pipeline for the Splunk and application fleet. Immutable, hardened images mean no configuration drift — every instance boots the same way and rebuilds cleanly.
- Automate infrastructure lifecycle in Jenkins Job DSL (Groovy), including S3 and SSM teardown as code with scoped least-privilege permissions. Also upgraded the Terraform runner image (**0.12 → 0.13**) without breaking the pipeline.
- Manage ECS capacity and scaling for the platform. Raised the production cluster's max task count to **150** after checking the underlying capacity, giving autoscaling more room under load.
- Led network security and IAM hardening: locked down the internal geocoding (nominatim / OSM) ELB to internal-only traffic across all **four** environments, cleaned up Lambda IAM roles by removing unused permissions, and managed CVE / vulnerability exceptions in the scanner.
- Drive cloud cost work across the Photos & Media estate. Three main levers: decommissioning stale ASGs, Launch Templates, and dead resources; Jenkins Job DSL teardown automation for S3 and SSM so cleanups are repeatable instead of one-off; and the media image-recompression pipeline (raw → recompressed in S3) for an ongoing storage-cost saving. All shipped as reviewed, reversible PRs.
- Consolidated the EC2 fleet onto Launch Templates. Built new Auto Scaling Groups on Launch Templates and removed the older Launch Configurations / legacy LTs in Terraform. No downtime, rolled out one environment at a time (dev, beta, stage, prod).
- Run observability and on-call. I built and run the Splunk logging platform, **own VividCortex database monitoring** for our RDS fleet (slow-query and query-plan analysis to catch regressions early), and work with SignalFX, RUM, synthetics, and APM across the photo-upload and CreateMedia paths. During incidents I use fast revert PRs as the main rollback. **24×7** on-call rotation.
- Handled several high-severity production incidents as primary on-call. Examples: a media-cloud traffic surge of **4.4M+** calls that overloaded the user-service Redis cluster; a **45-minute** Sev-1 outage caused by a Terraform / Akamai misconfiguration (fixed with a revert PR); and a **1-hour** platform-wide upload failure triggered by a third-party ingestion API.
- Cut on-call alert fatigue by about **25%** across the Photos, Media, and Upload dashboard groups. Did a domain-wide SignalFX cleanup — retired **147** old detectors and **617** stale charts so triage is faster and we get fewer noise pages.
- Lead RCAs and post-incident follow-ups for critical failures. Drove long-term fixes like enforcing timeouts on social-media ingestion calls and fixing autoscaling instability that came up in load tests, so the same problems don't come back.
- Support the media and AI platform at the infrastructure layer. I own the S3 bucket policy for the media-image-recompressor (raw → recompressed pipeline), the ECS capacity the media / AI services scale on, and AI config via SSM parameters. This is MLOps as the reliability layer, not model training.

KEY PROJECTS

- **Splunk Logging Platform (built from scratch):** Set up the license master and heavy forwarders entirely in Terraform (Launch Templates, ASGs, AMIs). A self-healing, code-managed logging layer the Photos & Media team uses every day for troubleshooting.
- **Cloud Cost Optimisation: Three** levers — decommissioning stale ASGs, Launch Templates, and dead resources; Jenkins Job DSL teardown as code for S3 and SSM so cleanups are automated; and the media image-recompression pipeline (raw → recompressed in S3) for ongoing storage savings. All shipped as reviewed, reversible PRs.
- **Observability & Incident Response:** Built and run Splunk. Own VividCortex database monitoring (slow-query and query-plan analysis on RDS). Use SignalFX, RUM, APM, and synthetics across the photo-upload and CreateMedia KPIs. Fast revert PRs are my main rollback during incidents.

EDUCATION

Master of Technology (M.Tech) — Indian Institute of Technology Kharagpur (IIT Kharagpur) | 2020 – 2022
Bachelor of Technology (B.Tech) — University of Petroleum and Energy Studies (UPES), Dehradun, India | 2015 – 2019